

Unsupervised Learning and Dimensionality Reduction

March 28, 2015

1 Unsupervised Learning and Dimensionality Reduction

```
In [455]: %matplotlib inline
          from algo_evaluation.handle_imports import *

In [132]: %%javascript
          IPython.load_extensions('calico-spell-check')

<IPython.core.display.Javascript object>
```

The goal is for you to think about how these algorithms are the same as, different from, and interact with your earlier work.

1.0.1 Classification Data Review

Higgs Data quick review:

- given outcomes of particle decays, detect Higgs boson;
- most of the **supervised** algorithms gave acceptable accuracy ranging from **0.8 - 0.9** far outperforming the **backpropagation** algorithm for neural networks (**0.5**)
- **learning weights** with Randomized Optimization algorithms improved accuracy of the neural network to **0.7** and higher

Additional Changes:

- features were rescaled as a requirement to feature transformation algorithms
- previously manually pruned features were put back to allow feature selection algorithms to autotically choose the most informative ones
- dataset was not reduced in size since dimensionality reduction provides speed advantage (the main reason for sampling records in the previous analysis)

```
In [374]: higgs_data = datasets.load_higgs_train(sample_size=None, verbose=True, scale=True,
          all_higgs_features = higgs_data[0].shape[1]
          prune_features=False)
```

Size of the dataset: 68114

Number of features: 30

Number of positives (signal): 31894

Number of negatives (background): 36220

Converters Data quick review:

- given online user behavior and preferences predict whether or not ad display will result in conversion;
- in contrast to Higgs dataset, supervised algorithm all provided high accuracy >94%
- this dataset was not subjected to randomized optimization analysis

```
In [392]: bid_data = datasets.load_bidding_train(verbose=True, scale=True)
          all_bid_data_features = bid_data[0].shape[1]
```

```
Size of the dataset: 57970
Number of features: 15
Number of converters: 399
Number of non-converters: 54615
Number of leads: 2956
```

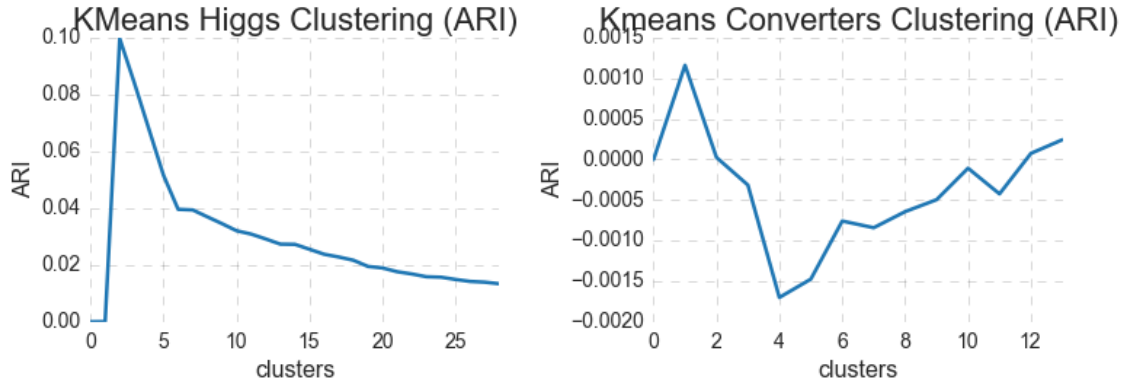
1.0.2 K-means Clustering

Evaluating cluster performance is different from the techniques used in Supervised Algorithms where accuracy was estimated roughly as a count of wrong answers. For clustering we must employ similarity metric: did clustering define separations of data similar to ground truth. In this analysis I shall use **Adjusted Rand Index** that measures the similarity of the two assignments, ignoring permutations and with chance normalization.

```
In [350]: df_kmeans_higgs = kmeans_eval.estimate_clusters(higgs_data)
```

```
In [367]: df_kmeans_converters = kmeans_eval.estimate_clusters(bid_data)
```

```
In [368]: plot_cluster_estimation.plot_kmeans_cluster_score(df_kmeans_higgs, df_kmeans_converters)
```



Choosing K:

Best data separation for **Higgs** data achieved with number of clusters equal **2** which makes sense since there are two classes in the data: signal and background, so the clustering lined up with labels. As number of the clusters grow, similarity measure gradually approaches zero.

Converters dataset achieves best separation for a single class which is the same as saying that clustering algorithm could not split data into groups. Since a single cluster does not make sense, it is reasonable to analyze next best separation. There is no clear convergence in this case towards a certain number of clusters and clearly cluster assignment does not line up with labels, so it is ok to choose an arbitrary number (we can stick with **10** and see how well that compares with dimensionality reduction algorithms)

Note, however, that neither of the above data separation is considered good (ARI=1.0 is a perfect score). Since scores are close to 0.0, label assignment is rather uniform especially for converters data. This suggests that datasets have more complex structure and are not easily separable into groups.


1.0.3 Expectation Maximization

As a second approach to clustering we will use the probabilistic model with assumption that all data points are generated from the number of Gaussian distributions. As a measure metric, we can use the scoring method which compute the **log probability** under the model.

```
In [355]: df_gmm_higgs = gmm_eval.estimate_clusters(higgs_data)
```

```
In [369]: df_gmm_converters = gmm_eval.estimate_clusters(bid_data)
```

```
In [370]: plot_cluster_estimation.plot_gmm_cluster_score(df_gmm_higgs, df_gmm_converters)
```



For **Higgs** dataset, loglikelihood of data is increasing as number of latent variables (here number of gaussian distributions) is increasing which suggests that every single feature came from a gaussian distribution with unique parameters (means and variances).

For **Converters** dataset, loglikelihood is increasing as well, however it is important to note that in the range of clusters from 4 to 9 probability increase is not very significant, so it is ok to make an assumption that data generally comes from 4 distinct distributions.

1.0.4 Feature Selection and Transformation for Higgs Dataset


In the previous analysis Higgs records were sampled and features were manually pruned due to the slow speed of supervised algorithms. In this analysis, complete dataset is restored to take advantage of automatic feature selection using dimensionality reduction algorithms.

Dimensionality reduction is the task of deriving a set of features that is smaller than the original feature set while retaining most of the variance of the original data. To estimate the best number of components to be used for data transformation, I used **data variance** and **reconstruction error** analysis to choose optimal setting.

```
In [162]: pca_rank = pca_eval.rank_features(higgs_data, n_components=all_higgs_features)
```

Using **PCA**, I extracted variances for each feature and sorted them from largest to smallest. Visualizing the results made it clear that “Derived” features from Higgs dataset have larger spread than “Primitives” and thus are more informative for classification. This supports the decision made in the previous analysis where “Primitive” features were pruned based on feature by feature analysis.

```
In [282]: pca_eval.plot_rank(pca_rank, title='Higgs Feature Rank')
```



Reducing features while retaining data variance:

When applying dimensionality reduction, it is important to retain certain threshold of data variance. In addition to visual inspeaction, I selected features which retained 95% of data variance (dimensionality reduced from 30 to 7). Algorithm ranked DER_mass_MMC the highest which makes sense since this feature estimates mass of the Higgs boson candidate and indeed is the most informative.

```
In [168]: reduced_features = pca_rank.variance_ratio[pca_rank.variance_ratio > 0.05]
reduced_features
```


```
Out[168]: feature
DER_mass_MMC                0.207485
DER_mass_transverse_met_lep  0.112727
DER_mass_vis                 0.087133
DER_pt_h                     0.058547
DER_deltaeta_jet_jet         0.058122
DER_mass_jet_jet             0.057542
DER_prodelta_jet_jet         0.053497
Name: variance_ratio, dtype: float64
```

Finding reconstrcution error: Another method for estimating number of components is to analyze reconstruction error and choose dimensionality with acceptable thresholds (in this case 30% error).

Note: all algorithms gave the same reconstruction error so it is sufficient to visualize results of one of them (for example, ICA)

```
In [195]: ica_comp = ica_eval.estimate_components(higgs_data)

In [230]: reconstruction_error = ica_comp.plot(x='components',
                                              y='reconstruction_error',
                                              kind='scatter',
                                              figsize=(8,4), title='Reconstruction Error (ICA)')
```



```
In [286]: reduced_higgs_dimension = df[df.reconstruction_error < 30].components.min()
          print 'Estimated number of reduced components for Higgs dataset =', reduced_dimension
```

Estimated number of reduced components for Higgs dataset = 7


With both methods for estimating number of components, 7 seems to be the optimal and thus shall be considered for dimensionality reduction for further analysis of the Higgs data.

1.0.5 Feature Selection and Transformation for Converters Dataset

Same principles and technique were applied to select features for converters dataset (retaining data variance of 95%)

```
In [251]: pca_rank_bid = pca_eval.rank_features(bid_data, n_components=all_bid_data_features)
```

```
In [284]: pca_eval.plot_rank(pca_rank_bid, title='Converters Feature Rank', figsize=(10,3))
```



```
In [300]: red_conv_feat = pca_rank_bid.variance_ratio[pca_rank_bid.variance_ratio > 0.05]
          reduced_conv_dimension = len(red_conv_feat)
          print 'Estimated number of reduced components for Converters dataset =', reduced_conv_dimension
          print red_conv_feat
```

Estimated number of reduced components for Converters dataset = 10

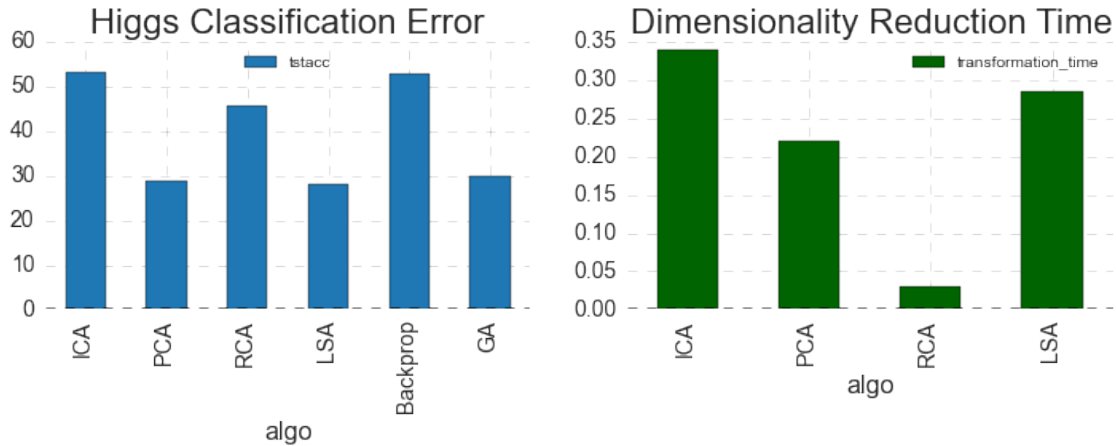
```
feature
source_timestamp    0.168507
user_local_hour     0.107505
user_id             0.100265
ipaddress           0.097562
country             0.085957
region              0.080283
metrocode           0.076695
postal_code         0.069630
week_day            0.062470
creative_uid        0.054318
Name: variance_ratio, dtype: float64
```

1.0.6 Neural Networks post Dimensionality Reduction

In the following experiment setup Higgs data set is reduced using four algorithms: PCA, ICA, RCA and LSA. Same error measure is applied as in the previous analysis where Higgs boson was classified with Neural Networks.

```
In [408]: df_nn = nn.evaluate_nn_accuracy(higgs_data, reduced_higgs_dimension)
```

```
In [468]: nn.plot_evaluation(df_nn)
```



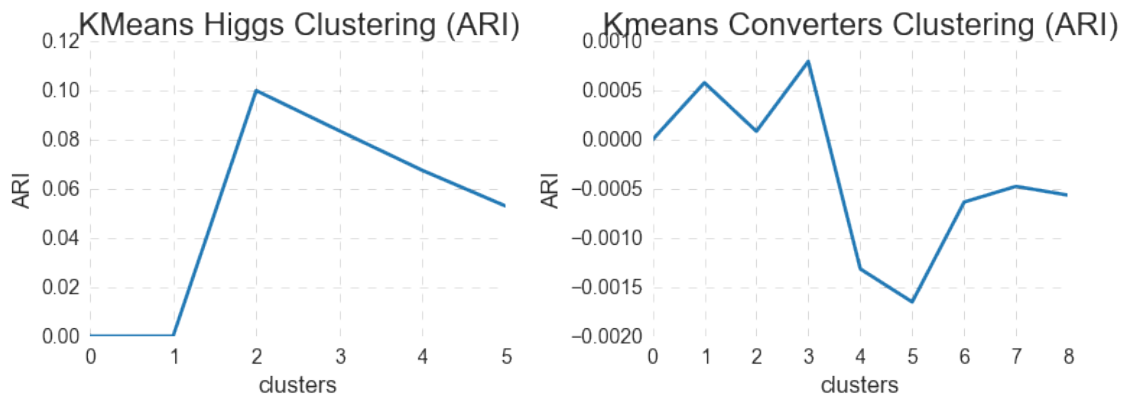
In addition to plotting the classification error post dimensionality reduction, I appended results from another two algorithms (Backpropagation and Weights Learning with Genetic Algorithm) to be used as a baseline. Very notable is the fact that PCA dimension reduction has improved classification error as compared to backpropagation and is on par with Genetic Algorithm which performed best for learning weights. Reducing dimensions with randomized projection does have not the best accuracy, however the speed of the algorithm is amazingly fast compared to all others and it can be very beneficial in the cases with large data dimensionality where accuracy could be sacrificed.

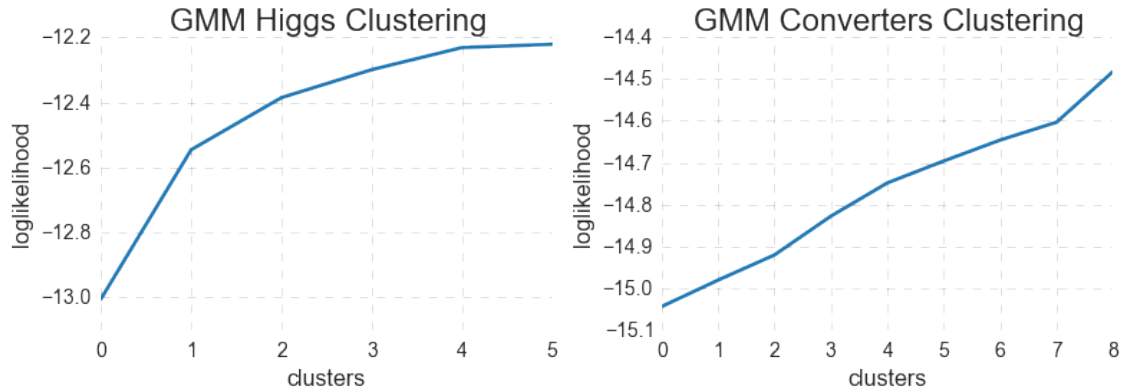
1.0.7 Clusteing post Dimensionality Reduction

In [464]: `df_higgs_dim = clustering_dim_reduction.evaluate_higgs_clustering(higgs_data, reduced_higgs_dim)`

In [476]: `df_conv_dim = clustering_dim_reduction.evaluate_conv_clustering(bid_data, reduced_conv_dimensions)`

In [486]: `clustering_dim_reduction.plot_cluster_performance(df_higgs_dim, df_conv_dim)`





Apply the clustering algorithms to the same dataset to which you just applied the dimensionality reduction algorithms (you’ve probably already done this), treating the clusters as if they were new features. In other words, treat the clustering algorithms as if they were dimensionality reduction algorithms. Again, rerun your neural network learner on the newly projected data.

If you used data that already had labels (for example data from a classification problem from assignment #1) did the clusters line up with the labels? Do they otherwise line up naturally? Why or why not? Compare and contrast the different algorithms. What sort of changes might you make to each of those algorithms to improve performance? How much performance was due to the problems you chose? Be creative and think of as many questions you can, and as many answers as you can. Take care to justify your analysis with data explicitly. Can you describe how the data look in the new spaces you created with the various algorithms? For PCA, what is the distribution of eigenvalues? For ICA, how kurtotic are the distributions? Do the projection axes for ICA seem to capture anything “meaningful”? Assuming you only generate k projections (i.e., you do dimensionality reduction), how well is the data reconstructed by the randomized projections? PCA? How much variation did you get when you re-ran your RP several times (I know I don’t have to mention that you might want to run RP many times to see what happens, but I hope you forgive me)? When you reproduced your clustering experiments on the datasets projected onto the new spaces created by ICA, PCA and RP, did you get the same clusters as before? Different clusters? Why? Why not? When you re-ran your neural network algorithms were there any differences in performance? Speed? Anything at all? It might be difficult to generate the same kinds of graphs for this part of the assignment as you did before; however, you should come up with some way to describe the kinds of clusters you get. If you can do that visually all the better.

In []: